

Extended Abstracts of the CHI Conference on Human Factors in Computing Systems (CHI EA '25) | Yokohama, Japan, April 26—May 01, 2025

Context over Categories: Implementing the Theory of Constructed Emotion with LLM-Guided User Analysis

TU Wien Informatics
Institute of Information Systems Engineering
Research Unit: Data Science

Nils Klüwer, Dr. Irina Nalis-Neuner, Assistant Prof. Dr. Julia Neidhardt

Overview & Contributions

Problem: Traditional emotion analysis relies on oversimplified categories (e.g., Ekman's universal emotions), neglecting context and individual differences.

Our Approach: We operationalize Lisa Feldman Barrett's Theory of Constructed Emotion, which emphasizes the crucial role of context in shaping emotions.

Key Contributions:

- Novel Framework: A context-aware emotion analysis pipeline using advanced LLM guidance.
- "Context Sphere": A new data construct representing a user's complete online behavior and interactions.
- LLM Adaptation: Demonstrates how to use LLMs for nuanced, context-sensitive emotion analysis.
- Focus: Content moderation and personalized systems.

Research Question

How can we leverage Large Language Models (LLMs) to analyze emotions in a way that reflects their complex, constructed nature, rather than relying on predefined categories?

Methodology

A. Context Sphere Building

Data Source: "Der Standard" online newspaper comments and related articles (May 2019, 23,925 users).

Goal: Create a comprehensive "context sphere" for each user, encapsulating:

- All comments made by the user.
- Surrounding context (article metadata, conversation threads).
- Interactions with other users (replies, etc.).

Key Features:

- Selective Pruning:** Includes relevant conversation threads up to the user's last comment, balancing context with efficiency.
- Privacy-Preserving:** Excludes usernames and gender to minimize bias.
- Markdown Format:** Optimized for readability by both humans and LLMs.

B: LLM-Guidance

Core Idea: Combine the Theory of Constructed Emotion with advanced LLM techniques.

Key Techniques:

- Role-Play Prompting:** The LLM (Gemini-1.5 Pro) impersonates Lisa Feldman Barrett. (Include a small, simplified version of Figure 2 if space allows).
- Controlled Generation:** A JSON schema guides the LLM's output, ensuring a structured analysis.
- Meta-Prompting:** A second LLM request generates a concise, human-readable Markdown report using a generalized template to adapt to specific instances.

Role-Setting Prompt	Role-Feedback Prompt	Instruction Prompt
User: You are Dr. Lisa Feldman Barrett, an expert in emotion analysis using your theory of constructed emotions. Your task is to ... [role setting]	Model: Understood. As Dr. Barrett, I will diligently analyze the provided text with a commitment to excellence. I will ... [anchoring the role]	User: Please analyze the following ... [Instruction] ...[context sphere]...

Figure 1: Role-Play Prompting according to the example from Kong et al. [28]

Output Structure: A multi-class JSON with five main fields the LLM must generate its response into:

- Core Affect Analysis
- Cognitive Appraisal & Conceptualization
- Cultural & Social Context
- Emotion Construction Analysis
- Emotional Dynamics & Changes

Each field includes sub-fields like "Thought Process," "Analysis," "Observable Patterns," "Anomalous Behaviour" and "Rationale."

Contact

Nils Klüwer



e12229263@student.tuwien.ac.at



Julia Neidhardt



recsys-lab.at/team/julia-neidhardt/

funded by:



Christian Doppler
Forschungsgesellschaft

Publications

- Kolb, T., Nalis-Neuner, I., & Neidhardt, J. (2023). Like a Skilled DJ—an Expert Study on News Recommendations Beyond Accuracy. In Proceedings of the International Workshop on News Recommendation and Analytics co-located with the 2023 ACM Conference on Recommender Systems (RecSys 2023). CEUR-WS.org.
- Thomas Elmar Kolb, Irina Nalis, Mete Sertkan, and Julia Neidhardt. 2022. The Role of Bias in News Recommendation in the Perception of the Covid-19 Pandemic. FAccTRec Workshop '22, in conjunction with RecSys 2022, Seattle, WA, USA.
- Kolb, T., E. Wagne, A. Sertkan, M., & Neidhardt, J. (2023). Potentials of Combining Local Knowledge and LLMs for Recommender Systems. In V. W. Anelli, P. Basile, G. De Melo, F. Donini, A. Ferrara, C. Musto, F. Narducci, A. Ragone, & M. Zanker (Eds.), Proceedings of the Fifth Knowledge-aware and Conversational Recommender Systems Workshop co-located with 17th ACM Conference on Recommender Systems (RecSys 2023) (pp. 61–64). CEUR-WS.org. <https://doi.org/10.34726/5334>
- Godo, L., & Neidhardt, J. (2024). Unlocking the Potential of Content-based Restaurant Recommender Systems. In Information and Communication Technologies in Tourism 2024 (pp. 239–244). Springer Nature Switzerland.
- Huebner, B., Kolb, T., & Neidhardt, J. (2024). Evaluating Group Fairness in News Recommendations: A Comparative Study of Algorithms and Metrics. In Adjunct Proceedings of the 32nd ACM Conference on User Modeling, Adaptation and Personalization (UMAP Adjunct '24) (pp. 10). ACM.
- Scholz, F., Kolb, T., & Neidhardt, J. (2024). Classifying User Roles in Online News Forums: A Model for User Interaction and Behavior Analysis. In Adjunct Proceedings of the 32nd ACM Conference on User Modeling, Adaptation and Personalization (UMAP Adjunct '24) (pp. 10). ACM.

image: "Flaticon.com". This poster has been designed using resources from Flaticon.com




Figure 2: Pipeline from Preprocessing to Final Output - Please note: User Study not part of this Late-Breaking Work Submission.

Methodology

C: Evaluation

Uses LLM-as-a-Judge (GPT-4o, Claude 3.5, Gemini 1.5 Flash) to assess LLM output

- Detecting **confabulations** (hallucinations) – confident but misleading outputs.
- Enables rapid, iterative development and refinement of the LLM pipeline.

Cost: Initial JSON analysis: ~\$0.035 per user. Full pipeline (including Markdown report): ~\$0.072 per user. Complete pipeline, including confabulation checks: ~\$0.462

Upcoming User Study not part of this Late Breaking Work

Table 1: Examples output from first LLM request

Field	Sub-Field	Output
Core Affect Analysis	Arousal	Generally high, fluctuating between agitated and calmly contemptuous. [...]
Cognitive Appraisal & Conceptualization	Analysis	[...] The user's interpretations are often filtered through a lens of skepticism and cynicism, particularly towards opposing political views. [...]
Cultural & Social Context	Thought Process	I will examine the cultural and social context by considering the platform ("Der Standard"), the political climate of Austria in May 2019 (pre-election period), and the specific topics discussed (immigration, politics, media). [...]
Emotion Construction Analysis	Rationale	[...] Their negative emotional responses are often amplified by their pre-existing beliefs and the context of the online forum. The occasional deviations from this pattern [...]
Emotional Dynamics & Changes	Anomalous Observations	The user's positive comment about Fendrich deviates from their usual negative pattern, suggesting a momentary shift in emotional state. [...]

Conclusion

This work advances emotion analysis by bridging cognitive science and computational systems, paving the way for nuanced, context-aware applications like content moderation and personalization. By operationalizing Barrett's Theory of Constructed Emotion \citet{barrett_context_2022} using the "context sphere" and LLMs, we show how dynamic emotion insights can be achieved. As a Late-Breaking Work, we highlight future exploration, including refining the framework and conducting user studies. The advancements have significant implications for developing more accurate and effective human-centered technologies. This approach captures the complexity of human emotions far better than models rooted in Ekman's universal emotions \citet{ekman_constants_1971}. These contributions address ethical risks highlighted by Crawford \citet{crawford_atlas_2021} and align with responsible AI initiatives \citet{werthern_digital_2024}.

Limitations and Future Work

This work acknowledges limitations in operationalizing the Theory of Constructed Emotion with current LLMs. The dynamic nature of emotions, as defined by Barrett, poses a challenge for evaluation, lacking a definitive "ground truth" \citet{barrett_navigating_2021}. Our current LLM-as-judge evaluation \citet{zheng_judging_2023}, while pragmatic, lacks human validation. Translating Barrett's framework required balancing theoretical flexibility with computational structure. The "context sphere" is a design choice to capture user context, but consistent analysis across diverse users remains a challenge. Future work will address these limitations through user studies, comparing LLM analyses to human interpretation. Scalability and handling potential failure modes, such as overgeneralization or excessive granularity, are also critical areas for future development.